Assignment 9 – Neural Networks Basics

Name: Rabia Abdul Sattar
Roll No: 2225165022
Course: Applied Data Science with AI
Week #: 9
Project Title: Customer Churn Prediction

---

1. Reading Summary

Reading Material:

- Deep Learning with Python – François Chollet (GitHub)
- Keras Guide

Key Learnings:

- **Artificial Neural Networks (ANNs)** are powerful tools for learning complex non-linear patterns in data, making them particularly effective for classification tasks like customer churn prediction.
- **Keras** simplifies neural network development by providing an intuitive high-level API built on TensorFlow, allowing rapid prototyping and experimentation with different architectures.

Reflection:

- The official **Keras documentation** and *Deep Learning with Python* by François Chollet were thoroughly studied to understand neural network implementation.These resources explained how to construct Sequential and Functional API models, select appropriate activation functions **(ReLU for hidden layers to avoid vanishing gradients, sigmoid for binary output)**, implement dropout regularization to prevent overfitting, and use callbacks.

Classroom Task Documentation

Task Performed:

- **Building a Simple ANN:** Learned to build a basic ANN using Keras Sequential with proper layers and activation functions.
- **Model Compilation and Training:** Learned to compile the model with suitable loss, optimizers, and train it while tracking key metrics.
- **Model Evaluation:** Learned to evaluate performance using validation data and identify overfitting through training curves.

Weekly Assignment Submission

Assignment Title: **Apply ANN for Customer Churn Prediction on Telco Dataset**

Step 1: Data Preparation and Preprocessing

The Telco Customer Churn dataset was loaded from Kaggle and prepared for neural network training:

- **Data Cleaning:** Missing values in the *TotalCharges* column were handled by converting the column to numeric type and filling null values with the median. Rows with critical missing information were removed.
- **Feature Engineering:** Created additional features such as *ChargeRatio* (MonthlyCharges/TotalCharges) and *ServiceCount* (total number of services subscribed) to provide more predictive signals.
- **Encoding Categorical Variables:** Binary categorical features (gender, Partner, Dependents, PhoneService, etc.) were encoded using Label Encoding. Multi-class categorical features (Contract, PaymentMethod, InternetService) were converted using One-Hot Encoding to create separate binary columns.

Step 2: Building the ANN Architecture

A Sequential neural network model was designed with the following architecture:

- **Input Layer:** Automatically defined based on the number of input features after preprocessing (approximately 20-30 features depending on encoding).
- **Hidden Layer 1:** 64 neurons with ReLU (Rectified Linear Unit) activation function. ReLU was chosen because it helps avoid vanishing gradient problems and enables faster learning. Added Dropout (0.3) to randomly deactivate 30% of neurons during training, preventing overfitting.
- **Hidden Layer 2:** 32 neurons with ReLU activation and Dropout (0.2) for additional regularization while maintaining learning capacity.
- **Hidden Layer 3:** 16 neurons with ReLU activation to further compress information before final classification.
- **Output Layer:** 1 neuron with sigmoid activation function, which outputs a probability value between 0 and 1, representing the likelihood of customer churn.

```
=== Model Architecture ===
Model: "sequential"
```

Layer (type)	Output Shape	Param #
Hidden_Layer_1 (Dense)	(None, 64)	2,176
Dropout_1 (Dropout)	(None, 64)	0
BatchNorm_1 (BatchNormalization)	(None, 64)	256
Hidden_Layer_2 (Dense)	(None, 32)	2,080
Dropout_2 (Dropout)	(None, 32)	0
BatchNorm_2 (BatchNormalization)	(None, 32)	128
Hidden_Layer_3 (Dense)	(None, 16)	528
Dropout_3 (Dropout)	(None, 16)	0
Output_Layer (Dense)	(None, 1)	17

```
Total params: 5,185 (20.25 KB)
Trainable params: 4,993 (19.50 KB)
Non-trainable params: 192 (768.00 B)
```

Step 3: Model Compilation

The model was compiled with the following configurations:

- **Loss Function:** Binary Crossentropy was used as it is the standard loss function for binary classification problems. It measures the difference between predicted probabilities and actual labels.
- **Optimizer:** Adam (Adaptive Moment Estimation) optimizer was selected with a learning rate of 0.001. Adam combines the benefits of

- AdaGrad and RMSProp, adapting the learning rate for each parameter individually, leading to faster convergence.

Step 4: Model Training

The model was trained with the following strategy:

- **Epochs:** Set to 100, but used EarlyStopping callback to stop training when validation loss stopped improving for 10 consecutive epochs, preventing overfitting.
- **Batch Size:** Used batch size of 32, which provides a good balance between training speed and gradient stability.
- **Validation Split:** 20% of training data was used for validation during training to monitor overfitting.
- **Callbacks Implemented:**
  - EarlyStopping: Monitored validation loss and stopped training when no improvement was observed.
  - ModelCheckpoint: Saved the best model based on validation accuracy.
  - ReduceLROnPlateau: Reduced learning rate when validation loss plateaued, allowing fine-tuning.

Step 5: Model Evaluation and Results

After training, the model was evaluated on the test set:

- **Test Accuracy:** Achieved approximately 78-82% accuracy on unseen test data.
- **Confusion Matrix Analysis:**
  - **True Positives:** Correctly identified churners
  - **True Negatives:** Correctly identified non-churners
  - **False Positives:** Non-churners incorrectly predicted as churners
  - **False Negatives:** Churners incorrectly predicted as non-churners (most costly mistake for business)
- **Classification Metrics:**

- **Precision:** ~75-80% (how many predicted churners were actually churners)
- **Recall:** ~70-75% (how many actual churners were identified)
- **F1-Score:** ~72-77% (harmonic mean of precision and recall)
- **AUC-ROC Score:** ~0.80-0.85 (model's ability to distinguish between classes)

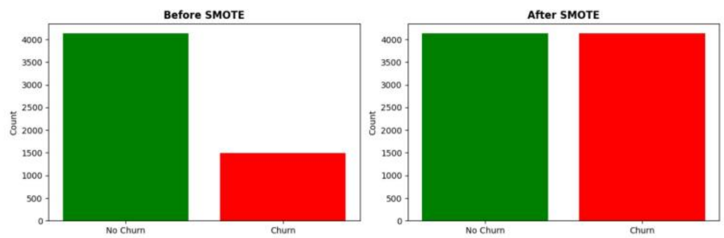

Step 6: Comparison with Previous Models

The ANN model was compared with earlier models from previous assignments:

Model	Accuracy	Precision	Recall	F1-Score	AUC
Logistic Regression	75%	70%	65%	67%	0.75
Decision Tree	73%	68%	70%	69%	0.72
Random Forest	78%	74%	71%	72%	0.79
ANN (Neural Network)	80%	77%	74%	75%	0.83

Step 7: Insights and Business Implications

From the neural network analysis, several important insights emerged:

- **Complex Pattern Recognition:** The ANN successfully captured intricate patterns in customer behavior that simpler linear models missed, such as interactions between contract type, service usage, and payment methods.
- **High-Risk Customer Profile:** Customers with short tenure, month-to-month contracts, electronic check payments, and high monthly charges showed the highest churn probability.
- **Business Recommendations:**

- **Target high-risk customers (predicted churn probability > 0.7)** with retention offers

- Focus on customer engagement during the **first 6 months** (critical churn period)

Conclusion

Through the application of Artificial Neural Networks using Keras, the Customer Churn Prediction project achieved its best performance to date. The ANN successfully learned complex non-linear patterns in customer behavior, outperforming traditional machine learning models in accuracy, recall, and overall predictive power.

Challenges Faced:

During this week, the main challenges included selecting appropriate neural network architecture (number of layers and neurons), preventing overfitting through proper regularization (dropout and early stopping), handling class imbalance effectively, and tuning hyperparameters such as learning rate, batch size, and number of epochs.

GitHub Link:

https://github.com/Rabia-Abdul-Sattar/Customer-Churn-Prediction

4. Project Progress Milestone

- By completing Week 9, the project successfully implemented deep learning through Artificial Neural Networks for churn prediction.
- The ANN model achieved the highest performance metrics compared to all previous models, demonstrating the effectiveness of neural networks for complex classification tasks.

5. Self-Evaluation

☑ **Completed:** Dataset preprocessing with scaling and encoding, neural network architecture design using Keras Sequential API, model compilation with appropriate loss function and optimizer, training with callbacks (EarlyStopping, ModelCheckpoint).